

CXO Agentic Skills Catalog

Introduction

Goal

Establish an initial CX Skills Catalog — the authoritative list of capabilities that Experience Cloud exposes to AI Harnesses/Orchestrators. These skills should be available across 1P and 3P agentic platforms (AEP AO/AI Assistant, Claude Code/Cowork, OpenAI Codex, and others). We aim to have a published catalog by Summit with as many skills as ready, plus a declared list of what’s coming post-Summit.

What Is a Skill?

A skill is a structured, composable unit of capability that an AI agent can discover, understand, and invoke — typically a markdown file that declares what the skill does, what inputs it needs, and when to use it. For a reference example, see github.com/anthropics/skills. The AEM team has already shipped 17 skills for Edge Delivery Services at github.com/adobe/skills. See [Learn More](#) below the catalog for details.

Open discussion points

On granularity: This catalog lists ~30 skills at a high level of abstraction, but depending on how teams decompose capabilities the real number could be hundreds or thousands. AEM EDS already has 17 for one product surface alone. The catalog will grow as teams structure their APIs for agentic consumption. This is an open discussion.

Caveats: Skills can be quite complex - almost software-level - so exposing them publicly when they are only a few weeks old may not be feasible. Once a skill is installed by a customer, the update cycle is not fully clear yet. Skills packaged for Codex or Cowork are very well-structured (see Figma or Cloudflare-deploy skills as examples) and almost always use MCP or very well-referenced APIs.

Skills Catalog (Draft for PM Review · March 2026)

Readiness Key: **A** = Available now · **B** = Targeted for Summit · **C** = Post-Summit · **P** = Placeholder (team engagement pending)

	Skill	Area	Product	Description	Sample Prompts	Readiness	Notes	PM Owner
1	Create an Audience (B2C)	Audiences B2C	RT-CDP	Build a new audience from a natural-language description. The system generates a segment definition, estimates size, and persists the audience after approval.	• "Build me an audience of first-time buyers who spent more than 5 minutes on product pages last week." • "Create an audience of loyalty members in California who haven't purchased in 90 days."	C		
2	Explore Audiences	Audiences B2C	RT-CDP	Search, filter, and inspect existing audiences — find sizes, check activation status, list destination mappings, and discover XDM fields available for targeting.	• "Which audiences haven't been activated in the last 30 days?" • "Show me fields I can use to target affluent buyers." • "What's my biggest audience?"	B		
3	Detect Audience Issues	Audiences B2C	RT-CDP	Find duplicate audiences, detect significant size changes, and surface anomalies across the audience inventory.	• "Do I have audiences with the same rules but different names?" • "Which audiences grew more than 20% this week?" • "What's my fastest-shrinking audience?"	P		
4	Create an Audience (B2B)	Audiences B2B	RT-CDP B2B / Marketo Engage / AJOB2B	Build a Smart List or person-based audience for B2B nurture and ABM programs — infer intent thresholds, persona coverage, and product-interest signals from a campaign brief.	• "Build an audience of mid-funnel contacts who attended the webinar and opened at least two emails."	C		Ambika Shetty Manasa Nagarajan
5	Create a Journey	Journeys B2C	AJO	Build a multi-step customer journey from a natural-language brief or visual reference — supports event-triggered, audience-targeted, business-event, and audience-qualification entry types across email, push, and SMS.	• "Create a post-purchase journey: send a thank-you push, wait 2 days, then branch by loyalty status." • "Build a re-engagement journey for dormant users with a 3-email drip." • "Here's a screenshot of the journey I want — build it."	A	Summit: Breakout + Lab. Text /ImageJourney on track (AJO26.3). Campaign create on track (AJO26.2). Image2Journey in beta for Summit. Deps: File upload (AO), Hybrid UI (AIA).	Ivan Mironchuk
6	Create a single step Action Campaign	Journeys B2C	AJO	Use natural language prompts to quickly create a scheduled single-action campaign.	"Create a campaign targeting newly registered users. Send a welcome email message with 10% discount offer. Schedule the campaign to run daily until 10/25."	B		Ivan Mironchuk
7	Create an Orchestrated Campaign	Journeys B2C	AJO	Use natural language prompts to create an orchestrated campaign with audience building and target attributes.				Frederic Mary

8	Create a Journey from Image	Journeys B2C	AJO	Use a reference image to quickly create a draft journey	Create a journey using the reference image	B		Ivan Mironchuk
9	Multi-turn edit a Journey	Journeys B2C	AJO	While creating a journey, ask for edits and changes.	Edit this journey, after the last push notification, wait 2 days and check for a reaction. If no reaction, send a follow-up push.	A		Ivan Mironchuk
10	Generate Journey Content	Journeys B2C	AJO	Create channel-specific content (email, push, SMS) for journey nodes using AI-powered generation, with template selection and tone control.	"Generate a welcome email for my onboarding journey with a friendly tone and 10% discount." "Regenerate the push notification with a more urgent tone."	A	Emagica content generation service on track for Summit (AJO26.2). Content create skill in MVP UX.	Amarnath Vannarath
11	Analyze a Journey	Journeys B2C	AJO	Identify drop-off points, detect audience overlaps across journeys, find scheduling conflicts, and surface operational metadata. ----- Identify journey anomalies, detect custom action errors, and proactively generate and display AI Recommendations.	"Where are users dropping off in my Holiday Campaign journey?" "Do any live journeys conflict with Summer Sale by schedule or audience?" "How many live journeys do I have?" ----- "Why am I seeing custom action errors in journey ABC?" "What is the health of journeyxyz?" "Why is there a drop in profiles in journeyXYZ?"	A ----- B	Summit: Breakout + Lab. Proactive monitoring (AI Recommendations) and anomaly detection on track (AJO26.3).	Jon Leininger
12	Simulate a Journey	Journeys B2C	AJO	Generate synthetic test profiles and events to validate a journey end-to-end before publication — detect broken logic, missing personalization, and content issues.	"Run a simulation across all paths of this journey." "Check for broken links and missing personalization before I publish."	B	Summit: Breakout + Lab. Auto-generate test data on track (AJO26.3). Reimagined in-product UX in progress. Deps: Context passing (AIA), Hybrid UI (AIA).	Nicolas DUFAU
13	Recommend Journey Strategy	Journeys B2C	AJO	Translate a business goal into recommended journey opportunities with evidence-backed justification — cohort traits, behavioral signals, and historical performance.	"I want to increase retention among high-value customers — what journey should I run?" "Which customer behaviors suggest an upsell opportunity right now?"	C	Journey Cohort Discovery in Alpha. Roadmap mention for Summit. Goal definition, goal tracking, and business context upload are open dependencies.	Gael Mouello
14	Create a Journey (B2B)	Journeys B2B	AJOB2B	Build person-based email nurture journeys from a campaign brief or Marketo engagement program — set cadence, map content, and select or create audiences for entry.	"Create an email nurture journey for leads in Project Falcon with weekly cadence." "Translate my Marketo engagement program into a Falco journey."	B		Ambika Shetty John Dictson
15	Research an Account	Accounts B2B	RT-CDP B2B / AJOB2B	Surface detailed account and prospect profiles — recent investments, strategic priorities, pain points, and engagement history.	"Give me a detailed profile for Acme Corp." "What are XYZ's recent strategic priorities and pain points?"			?
16	Draft Outreach	Accounts B2B	Sales Qualifier / AJOB2B / Marketo Engage	Generate personalized outreach emails and engagement sequences based on account context, recent engagement signals, and strategic goals.	"Draft an outreach email for ABC leveraging their recent engagement." "Change the opening to reference their latest earnings call."		Assume this is Sales Qualifier.	Tarun Arora
17	Manage Buying Groups	Accounts B2B	AJOB2B	Infer B2B buying-group composition from product intent signals, trace account paths through journeys, and recommend buying-group structures.	"Which roles are commonly represented for this account?" "Trace the most common path Decision Makers took in this journey."	A		Ambika Shetty Manasa Nagarajan
18	Visualize Data	Data & Insights	CJA	Answer natural-language questions about your analytics data by building visualizations (line, bar, donut, table, summary number) in Analysis Workspace. Designed for analysts and non-analysts alike, with guided onboarding and follow-up questions that build on prior context.	"Trend orders in July." "Show revenue by region." "What were my top 10 SKUs by profit last month?"			
19	Diagnose Root Causes	Data & Insights	CJA	Identify what's driving a metric change — surface contributing dimensions, quantify impact, and recommend areas to investigate further.	"Why did conversion drop last week?" "What's driving the revenue increase in EMEA?" "Which channels contributed most to the sign-up spike?"			
20	Analyze Experiments	Data & Insights	Experimentation Accelerator	Summarize experiment results, explain why treatments won or lost, identify content attributes driving performance, and recommend next experiments.	"What did we learn from this experiment?" "Why did treatment A outperform?" "What should I test next?"			
21	Onboard Data	Data Engineering	AEP	Analyze datasets, auto-generate schemas, run Semantic Enrichment and provide additional context on the fields in the data	"Create a schema for this data" "Analyze this file and tell me which fields can be ignored" "Run semantic enrichment on this data"	C		
22	Transform Data	Data Engineering		Translate natural language to SQL for data prep	"Write a query that joins my loyalty and web datasets, groups by region, and computes monthly revenue." "Optimize this SQL query" "Run this computation every month"	C		

23	Analyze data collection configurations	Data Engineering		Audit data collection configurations and diagnose the environment	"Audit my site's data collection tags and suggest optimizations."	C		
24	Troubleshoot Product Issues & Create Support Tickets	Support	AEP (cross-product)	Get step-by-step guidance to diagnose platform issues, with answers sourced from Adobe documentation and contextualized to your sandbox. Includes interactive diagnostics that run automated investigation routines against live sandbox data, support case creation and tracking, and context-aware escalation routing.	• "Why are profiles not qualifying for my audience?" • "Why are fewer profiles exported to my Facebook destination than expected?" • "Create a support ticket for my ingestion errors."	A (Interactive Diagnostics & Case Management available only for explorer program customers)		
25	Manage Data Lifecycle and Retention Policies on Datasets	Data Management	AEP	Manage dataset lifecycle and set up TTLs	"Help me evaluate my datasets and show me the ones which do not have TTL setup" "Which datasets are without an active dataflow" "Setup an expiration time (or TTL) of 30days on this event level dataset XYZ"	C		
26	Process Unstructured Data	Unstructured Data	AEP	Ability for customers to send unstructured data using existing connectors or upload directly into the AI assistant	• "Using these sales play documents, customer conversation calls and surveys docs, synthesize insights about the target account and come up with a good sales proposal for VP of customer engineering team" • "Determine the user's current intent." • "From the provided data, identify the buying group" • "Draft a personalized outreach email."	P		
27	Update Site Content	Content	AEM	Update, replace, or add content elements across AEM pages, content fragments, forms, and assets via natural-language instructions.	• "Replace the hero image on the homepage with the new campaign asset." • "Update the pricing section on the product page to reflect the March changes."	A		Corey Dulimba
28	Create Forms & Communications	Content	AEM	Build adaptive forms or generate personalized data-driven correspondence (statements, welcome kits, policy documents) from natural-language prompts.	• "Create a contact form with name, email, company, and a message field." • "Generate a welcome kit for new insurance policyholders."	A		Arun Taneja
29	Produce an Experience	Content	AEM	Create a new page or experience from a campaign brief and approved assets — the agent reads the brief, proposes a build plan, selects templates, sources content from the DAM, and assembles a draft for review.	• "Create a landing page for the Digital Leaders Webinar using the attached brief." • "Build a campaign page — here's the brief document."	B		Corey Dulimba
30	Migrate a Site	Content	AEM	Migrate non-AEM sites into AEM Edge Delivery Services — content, styles, blocks, and design — through conversational prompts.	• "Migrate this WordPress site to AEM Edge Delivery." • "Create a landing page block matching this design."	A		Gabriel Walt
31	Discover Content	Content	AEM	Search across AEM Assets, Content Fragments, and Forms to find the most relevant approved content for reuse — images, videos, PDFs, articles, templates.	• "Find me approved product images for the summer campaign." • "Show me all content fragments tagged 'financial services'."	A		Apoorva Gupta
32	Organize Content	Content	AEM	Create and manage DAM taxonomy — generate tag hierarchies, assign metadata, build folder structures, and maintain consistent classification across assets.	• "Create a tag hierarchy for our product lines." • "Auto-tag the new batch of campaign images." • "Propose a folder structure for the 2026 brand refresh assets."	B		Nick Whittenburg
33	Optimize Content	Content	AEM	Generate channel-ready asset variants — new renditions, background changes, visual property adjustments — from natural-language instructions.	• "Create a social-crop variant of this hero image." • "Adjust this product image for a dark-background email layout."	A		Greg Klebus
34	Check Brand Compliance	Content	AEM	Validate pages or assets against pre-ingested brand guidelines — tone, claims, logo usage, typography, imagery — and surface specific violations with remediation guidance. Supports ingesting brand URLs and running full-site scans.	• "Is this page aligned with my brand?" • "Check if the new landing page follows brand messaging guidelines." • "Scan our site against the latest brand policy."	A		Philippe Kapfer
35	Manage Permissions & Rights	Content	AEM	Enforce digital rights management and access control across assets — evaluate DRM rules, manage permission policies, and ensure content usage stays within licensing and governance boundaries.	• "Which assets in this campaign folder have usage restrictions?" • "Can I use this image in a paid social ad?" • "Show me assets expiring in the next 30 days."	A		Philippe Kapfer
36	Plan Campaigns & Briefs	Operations	Workfront	Generate Planning workspaces that organize campaigns, budgets, and connected records for end-to-end marketing ops visibility. Transform an idea or uploaded document into a structured, connected set of Planning records — taxonomies, timelines, and deliverables.	• "Set up a Planning workspace for the Q3 product launch." • "Here's our campaign idea doc — turn it into a structured brief with timeline and deliverables." • "Define taxonomies for our 2026 campaign calendar."			

37	Manage Projects & Resources	Operations	Workfront	Execute operational changes in natural language — create projects with task plans, bulk update ownership and status across work items, assign resources, adjust schedules and milestones, generate calculated field formulas, and track project health. Surface actionable insights from broad, exploratory queries across the Workfront ecosystem.	“Create a project for the spring launch with a full task plan.” “Reassign all of Maria’s open tasks to Jake.” “What tasks are overdue in the Q2 campaign?” “Write a calculated field formula for days until deadline.”			
38	Accelerate Content Reviews	Operations	Workfront	Manage multi-stage content review cycles — add, remove, or reorder review stages with deadlines, approvers, and reviewers. Automate brand compliance checks on assets during review and provide instant, actionable feedback. Send approved assets to AEM.	“Add a legal review stage before final approval on the summer campaign.” “Who still needs to approve the homepage redesign?” “Check these assets against brand guidelines and route approved ones to AEM.”			
39	Understand Data	Operations	Workfront	Surface actionable insights from Workfront Projects, Tasks, Issues across the Workfront workflow.	- Show me all active projects for the Brand Marketing team - What does the breakdown of my active tasks look like by status? - How many issues do we have in that Needs Triage status?	C		Matt Mitchell Oznur Aytekin
40	Content Generation & Variants	Content Production	GenStudio PeM	Generate on-brand marketing content variants across channels — emails, ads, social posts, banners — using approved brand templates and guidelines. (Placeholder — skills TBD with GenStudio PeM team.)	TBD			
41	Asset Generation	Content Production	Firefly	Generate and edit images, vectors, and design assets from natural-language prompts, with brand-safe controls and commercial licensing. (Placeholder — skills TBD with Firefly team.)	TBD			
42	Design Composition	Content Production	Express	Create and customize designs — social graphics, flyers, short videos — from templates or natural-language briefs. (Placeholder — skills TBD with Express team.)	TBD			
43	Video Review & Approval	Content Production	Frame.io	Manage video review workflows — collect feedback, manage versions, route approvals, and publish approved cuts. (Placeholder — skills TBD with Frame.io team.)	TBD			
44	Run an Experiment	Optimization & Experimentation	Target	Set up and launch an A/B test, multivariate test, or personalization activity from a natural-language brief or approved execution intent. Generates the activity configuration - variants, audiences, success metrics, traffic allocation - and deploys to the appropriate execution surface (Target or AJO Decisioning) after human approval. Can receive upstream intent from the Experimentation Accelerator via Agent Orchestrator.	“Set up an A/B test on the homepage hero with three headline variants targeting new visitors.” “Launch a personalization activity for returning customers with a 10% discount offer.” “Create an experiment based on the recommendation from the Experimentation Accelerator.”	B - aspirational	Mutation skill. AO-gated, human approval required before go-live. Complements EA’s “Analyze Experiments” (#16) which operates post-execution. Execution surface resolved by orchestrator based on channel and customer config.	
45	Run an Experiment	Optimization & Experimentation	AJO	Set up and launch an A/B test, multivariate test, or personalization activity from a natural-language brief or approved execution intent. Generates the activity configuration - variants, audiences, success metrics, traffic allocation - and deploys to the appropriate execution surface (Target or AJO Decisioning) after human approval. Can receive upstream intent from the Experimentation Accelerator via Agent Orchestrator.	“Set up an A/B test on the homepage hero with three headline variants targeting new visitors.” “Launch a personalization activity for returning customers with a 10% discount offer.” “Create an experiment based on the recommendation from the Experimentation Accelerator.”	C	Mutation skill. AO-gated, human approval required before go-live. Complements EA’s “Analyze Experiments” (#16) which operates post-execution. Experimentation Agent listed as Roadmap Mention for Summit (H1’26).	cprasad@adobe.com
46	Manage Experiment Lifecycle	Optimization & Experimentation	Target	Pause, resume, stop, schedule, or adjust running experiments - including traffic reallocation, priority changes, and bulk operations across campaigns. Supports single-activity and multi-activity mutations with blast-radius awareness.	“Pause the Black Friday homepage test.” “Stop all experiments in the Holiday campaign.” “Increase traffic allocation to variant B to 70%.” “Schedule the Spring Sale test to go live next Monday at 9am.”	B - aspirational	Mutation skill. Blast radius: Single or Multiple. Bulk operations require explicit preview-then-confirm. All state changes AO-gated.	
47	Manage Experiment Lifecycle	Optimization & Experimentation	AJO	Pause, resume, stop, schedule, or adjust running experiments - including traffic reallocation, priority changes, and bulk operations across campaigns. Supports single-activity and multi-activity mutations with blast-radius awareness.	“Pause the Black Friday homepage test.” “Stop all experiments in the Holiday campaign.” “Increase traffic allocation to variant B to 70%.” “Schedule the Spring Sale test to go live next Monday at 9am.”	C	Mutation skill. Blast radius: Single or Multiple. Bulk operations require explicit preview-then-confirm. All state changes AO-gated. Post-Summit (H1’26).	cprasad@adobe.com
48	Inspect Experiment State	Optimization & Experimentation	Target	Query what experiments are running, their configuration, targeting rules, change history, and delivery context across execution surfaces. Conversational interface to experiment state - answer natural-language questions about live operations. Read-only; no mutations, no insight synthesis.	“What experiments are currently running?” “What audience is targeted in the checkout A/B test?” “Has the traffic allocation changed since launch?” “What’s running in the homepage mbox?” “Which experiments are scheduled to start next week?”	B - aspirational	Read-only. No approval required. Distinct from EA’s “Analyze Experiments” (#16) which interprets results; this skill surfaces factual state only.	

49	Inspect Experiment State	Optimization & Experimentation	AJO Decisioning	Query what experiments are running, their configuration, targeting rules, change history, and delivery context across execution surfaces. Conversational interface to experiment state - answer natural-language questions about live operations. Read-only; no mutations, no insight synthesis.	• "What experiments are currently running?" • "What audience is targeted in the checkout A/B test?" • "Has the traffic allocation changed since launch?" • "What's running in the homepage mbox?" • "Which experiments are scheduled to start next week?"	C	Read-only. No approval required. Distinct from EA's "Analyze Experiments" (#16) which interprets results; this skill surfaces factual state only. Post-Summit.	cprasad@adobe.com
50	Retrieve Experiment Metrics	Optimization & Experimentation	Target	Surface raw performance data from running or completed experiments - lift, confidence intervals, conversion rates per variant, exposure counts, traffic volume, and statistical significance flags. Returns factual metrics only; does not interpret results, declare winners, or recommend actions.	• "What's the lift on the homepage test?" • "Show me conversion rates per variant for the Spring Sale experiment." • "How many users have been exposed to each variant?" • "Which experiments reached statistical significance last quarter?"	B - aspirational	Read-only. Explicitly complementary to "Analyze Experiments" (EA, #16): this skill retrieves raw data, EA interprets it. Separation ensures execution surface does not opine on strategy.	
51	Retrieve Experiment Metrics	Optimization & Experimentation	AJO	Surface raw performance data from running or completed experiments - lift, confidence intervals, conversion rates per variant, exposure counts, traffic volume, and statistical significance flags. Returns factual metrics only; does not interpret results, declare winners, or recommend actions.	• "What's the lift on the homepage test?" • "Show me conversion rates per variant for the Spring Sale experiment." • "How many users have been exposed to each variant?" • "Which experiments reached statistical significance last quarter?"	C	Read-only. Explicitly complementary to "Analyze Experiments" (EA, #16): this skill retrieves raw data, EA interprets it. Post-Summit.	cprasad@adobe.com
52	Manage Optimization Audiences & Offers	Optimization & Experimentation	Target	Create, modify, and apply audiences and offers scoped to experimentation and personalization activities. Delivery-level constructs distinct from RT-CDP platform audiences (#1-3). Supports creating targeting rules, defining offer content, and associating both with experiments or personalization activities.	• "Create a returning-customer audience for my personalization activity." • "Update the homepage offer headline to 'Spring Sale - 20% Off.'" • "What offers are associated with the checkout experiment?" • "Duplicate the mobile audience and add a geo filter for California."	B - aspirational		
53	Plan an Experiment	Data & Insights	Experimentation Accelerator	Design and validate an experiment before launch - calculate sample size and expected duration, determine how many treatments can be supported, forecast audience sizes, assess minimum detectable effect thresholds, and provide guidance on experimental design tradeoffs. Integrates with Audience Agent for sizing and uses traffic data from CJA/AA/Target to ground estimates.	• "How long will this A/B test take to reach stat sig with 3 variants?" • "Can my site support a 4-arm test on the checkout page?" • "What's the minimum lift I can detect with 2 weeks of traffic?" • "I want to run a test on mobile users - estimate the timeline."	B - Aspirational	Pre-experiment counterpart to "Analyze Experiments" (#16) which operates post-experiment. Outputs: required duration, supported arms, lift detection thresholds. Inputs include traffic volume, number of arms, projected lift, confidence level. Supports bidirectional calculation (duration-to-arms and arms-to-duration).	
54	Monitor Experiment Health	Data & Insights	Experimentation Accelerator	Continuously validate running experiments for statistical and implementation integrity. Detects sample ratio mismatch (SRM), pre-experiment bias between treatment groups, experiments unlikely to reach significance, missing data post-launch, and absent conversion tracking. Surfaces warnings proactively and provides guidance on remediation.	• "Are any of my running experiments showing SRM issues?" • "Is the checkout test likely to finish in time?" • "Has the new homepage test received data since it went live?" • "Check if any experiments have biased treatment groups."	C	In-flight monitoring - distinct from "Analyze Experiments" (#16) which interprets completed results, and from "Inspect Experiment State" (Target/AJO) which surfaces raw configuration. This skill applies statistical and implementation validation logic. Severity levels: critical (SRM, no data), warning (unlikely to complete, bias), informational (approaching completion).	
55	Generate Experiment Content	Data & Insights	Experimentation Accelerator	Generate AI-powered treatment variants - subject lines, marketing copy, and image suggestions - informed by experiment insights and opportunity analysis. Uses content attributes that drove past performance to recommend what to test next. Content is generated via eMagica /GenStudio and ranked by EA's attribute model to surface the top variants.	• "Generate three headline variants for the homepage based on what's worked before." • "What subject lines should I test next given the insights from my last email experiment?" • "Create treatment content for the opportunity card on the product page."	B - Aspirational	Insight-driven content generation, distinct from GenStudio PeM's general-purpose "Content Generation & Variants" (#32). EA generates content specifically informed by experiment insights and opportunity attributes. Content is ranked by EA's model before surfacing. Integrates with eMagica for generation and allows refinement via AI Assistant.	

Learn More

A skill is a self-contained folder with a markdown file that tells an agent what a capability does, when to use it, what inputs it needs, and what patterns to follow. Skills can be invoked via MCP, A2A, direct API, or a platform-native registry — the format is protocol-agnostic. They range from simple single-action skills (call one API endpoint) to complex orchestration skills that chain multiple tools together with reasoning logic.

For a concrete reference, see the Anthropic skills repository (github.com/anthropics/skills) which contains community-contributed skills for Claude Code — things like Figma integration, Cloudflare deployment, and code review. Each skill is a self-contained folder with a markdown description. That's the pattern.

Further reading:

- [Introducing Agent Skills](#) — Anthropic's official launch post (Oct 2025)
- [Skills Explained](#) — Anthropic engineering blog on progressive disclosure architecture, decision matrices for Skills vs Prompts/Subagents, and best practices
- [Simon Willison: "Claude Skills are awesome, maybe a bigger deal than MCP"](#) — independent technical deep-dive

- [Agent Skills documentation](#) — API reference and SDK integration